

```
In [1]: # Necessary Imports
import pandas as pd
import numpy as np

import seaborn as sns
import matplotlib.pyplot as plt

import statsmodels.api as sm

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #Import data
df = pd.read_csv('delivery_time.csv')
df.head(7)
```

```
Out[2]:
```

	Delivery_Time	Sorting_Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10
5	15.35	6
6	19.00	7

## EDA

```
In [3]: df.shape
```


```
Out[3]: (21, 2)
```

```
In [4]: df.skew()
```

```
Out[4]: Delivery_Time    0.352390
Sorting_Time    0.047115
dtype: float64
```

```
In [5]: #Co-relation of Target and Feature  
sns.heatmap(df.corr(), annot = True)
```

```
Out[5]: <AxesSubplot:>
```



```
In [6]: df.info() #no null values in dataset
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 21 entries, 0 to 20  
Data columns (total 2 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   Delivery_Time    21 non-null     float64  
1   Sorting_Time     21 non-null     int64  
dtypes: float64(1), int64(1)  
memory usage: 464.0 bytes
```

```
In [7]: df.describe()
```


```
Out[7]:
```

	Delivery_Time	Sorting_Time
count	21.000000	21.000000
mean	16.790952	6.190476
std	5.074901	2.542028
min	8.000000	2.000000
25%	13.500000	4.000000
50%	17.830000	6.000000
75%	19.750000	8.000000
max	29.000000	10.000000

## Graphical Univariate analysis on dataset


```
In [8]: sns.distplot(df.Delivery_Time)
```

```
Out[8]: <AxesSubplot:xlabel='Delivery_Time', ylabel='Density'>
```



```
In [9]: sns.distplot(df.Sorting_Time)
```

```
Out[9]: <AxesSubplot:xlabel='Sorting_Time', ylabel='Density'>
```



## Transforming data - Log scale

```
In [10]: df['dt'] = np.log(df['Delivery_Time'])  
df['st'] = np.log(df['Sorting_Time'])  
df.head()
```

```
Out[10]:
```

	Delivery_Time	Sorting_Time	dt	st
0	21.00	10	3.044522	2.302585
1	13.50	4	2.602690	1.386294
2	19.75	6	2.983153	1.791759
3	24.00	9	3.178054	2.197225
4	29.00	10	3.367296	2.302585


```
In [11]: df.skew()
```

```
Out[11]: Delivery_Time    0.352390  
Sorting_Time    0.047115  
dt             -0.451290  
st             -0.605236  
dtype: float64
```

## Graphical Bi-variate analysis on dataset


```
In [12]: #Pairplot of predictor and target  
sns.pairplot(df, vars = ['dt', 'st'])
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x23c241481c0>
```




```
In [13]: y = df['dt']
x = df['st']
plt.scatter(y,x)                                # Scatter plot of new Delivery_t

plt.xlabel('Sorting Time', fontsize = 10)        # Named the axes
plt.ylabel('Delivery Time', fontsize = 10)
plt.title(label='Scatter Plot', fontsize=20, y=1.05)
plt.show()                                     # Show the plot
```



## Fitting a Linear Regression Model

```
In [14]: sns.regplot(x="st", y="dt", data=df);
```



```
In [15]: import statsmodels.formula.api as smf

model = smf.ols('dt~st', data = df).fit()
model.summary()
```

Out[15]: OLS Regression Results

<b>Dep. Variable:</b>	dt	<b>R-squared:</b>	0.772
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.760
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	64.39
<b>Date:</b>	Fri, 21 Oct 2022	<b>Prob (F-statistic):</b>	1.60e-07
<b>Time:</b>	21:13:02	<b>Log-Likelihood:</b>	10.291
<b>No. Observations:</b>	21	<b>AIC:</b>	-16.58
<b>Df Residuals:</b>	19	<b>BIC:</b>	-14.49
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	1.7420	0.133	13.086	0.000	1.463	2.021
<b>st</b>	0.5975	0.074	8.024	0.000	0.442	0.753

<b>Omnibus:</b>	1.871	<b>Durbin-Watson:</b>	1.322
<b>Prob(Omnibus):</b>	0.392	<b>Jarque-Bera (JB):</b>	1.170
<b>Skew:</b>	0.577	<b>Prob(JB):</b>	0.557
<b>Kurtosis:</b>	2.916	<b>Cond. No.</b>	9.08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

'Sorting time' p value from summary = 0.000 Hence, the 'Sorting time' is an important feature for the Target 'Delivery Time'

The feature Const p value = 0.000 Hence, Constant is also an important feature for the Target 'Delivery Time'

R-squared is closer to 1 that means the regression model covers most part of the variance of the values of the response variable and can be termed as a good model.

```
In [16]: #Co-efficients values  
#beta1 and bet0 values  
model.params
```

```
Out[16]: Intercept    1.741987  
st              0.597522  
dtype: float64
```

```
In [17]: #t and p-Values for intercept and Sorting Time  
print(model.tvalues, '\n', model.pvalues)
```

```
Intercept    13.085552  
st           8.024484  
dtype: float64  
Intercept    5.921137e-11  
st           1.601539e-07  
dtype: float64
```