

```
In [1]: # Necessary imports
import pandas as pd
import numpy as np

import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv('Salary_Data.csv') #dataset
df.head(7)
```

```
Out[2]:
```

| | YearsExperience | Salary |
|---|-----------------|--------|
| 0 | 1.1 | 39343 |
| 1 | 1.3 | 46205 |
| 2 | 1.5 | 37731 |
| 3 | 2.0 | 43525 |
| 4 | 2.2 | 39891 |
| 5 | 2.9 | 56642 |
| 6 | 3.0 | 60150 |

EDA

```
In [3]: df.shape #shape of dataset
```

```
Out[3]: (30, 2)
```

```
In [4]: df.info() #no null value in dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 608.0 bytes
```

```
In [5]: df.skew() #skewness of Feature and Target
```

```
Out[5]: YearsExperience    0.37956
Salary          0.35412
dtype: float64
```

```
In [6]: #Co-relation of Target and Feature is close to 1
sns.heatmap(df.corr(), annot = True)
```

Out[6]: <AxesSubplot:>



In [7]: `df.describe()`


Out[7]:

| | YearsExperience | Salary |
|--------------|-----------------|---------------|
| count | 30.000000 | 30.000000 |
| mean | 5.313333 | 76003.000000 |
| std | 2.837888 | 27414.429785 |
| min | 1.100000 | 37731.000000 |
| 25% | 3.200000 | 56720.750000 |
| 50% | 4.700000 | 65237.000000 |
| 75% | 7.700000 | 100544.750000 |
| max | 10.500000 | 122391.000000 |

Graphical Univariate analysis on dataset


In [8]: `sns.distplot(df.YearsExperience)`

Out[8]: <AxesSubplot:xlabel='YearsExperience', ylabel='Density'>



```
In [9]: sns.distplot(df.Salary)
```

```
Out[9]: <AxesSubplot:xlabel='Salary', ylabel='Density'>
```



Transforming Target data - Log scale


```
In [10]: #using log for target data to scale target to feature range
df['log_salary'] = np.log(df.Salary) #new target= log_salary
df.head()
```

```
Out[10]:
```

| | YearsExperience | Salary | log_salary |
|---|-----------------|--------|------------|
| 0 | 1.1 | 39343 | 10.580073 |
| 1 | 1.3 | 46205 | 10.740843 |
| 2 | 1.5 | 37731 | 10.538237 |
| 3 | 2.0 | 43525 | 10.681091 |
| 4 | 2.2 | 39891 | 10.593906 |

```
In [11]: sns.distplot(df.log_salary)
```

Out[11]: <AxesSubplot:xlabel='log_salary', ylabel='Density'>



In [12]: `df['log_salary'].head()` *#scaled Target data scaled*

Out[12]:

| | |
|---|-----------|
| 0 | 10.580073 |
| 1 | 10.740843 |
| 2 | 10.538237 |
| 3 | 10.681091 |
| 4 | 10.593906 |

Name: log_salary, dtype: float64

In [13]: `df.skew()` *#skewness of target after scaling*

Out[13]:

| | |
|-----------------|-----------|
| YearsExperience | 0.379560 |
| Salary | 0.354120 |
| log_salary | -0.044126 |

dtype: float64


Graphical Bi-variate analysis on dataset

In [14]: *# Following the regression equation, our dependent variable (y) is the Salary*
`y = df['log_salary']`

Similarly, our independent variable (x) is the Year Experience
`x = df['YearsExperience']`

#Pairplot of predictor and target
`sns.pairplot(df, vars = ['log_salary', 'YearsExperience'])`

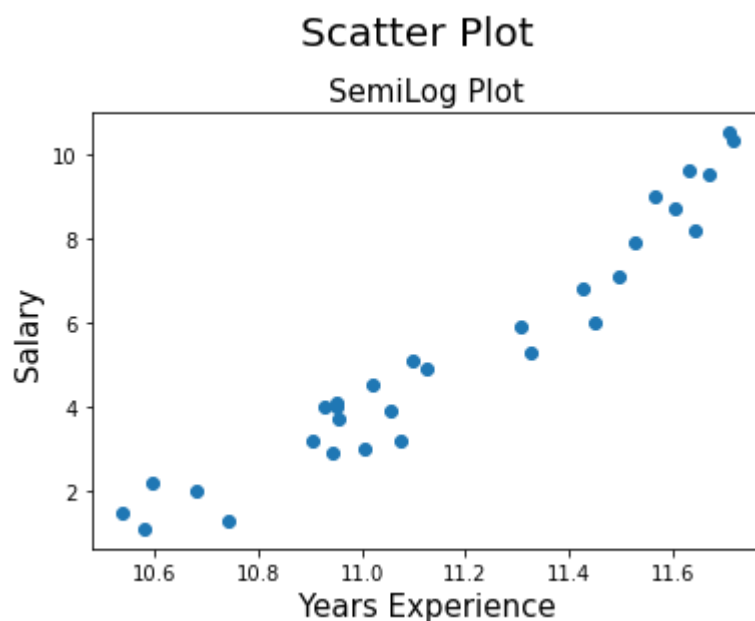
Out[14]: <seaborn.axisgrid.PairGrid at 0x271fa5469a0>



SemiLog Plot


```
In [15]: plt.scatter(y,x)                                     # Scatter plot

plt.xlabel('Years Experience', fontsize = 15)                 # Named the axes
plt.ylabel('Salary', fontsize = 15)
plt.title(label='SemiLog Plot', fontsize=15)
plt.suptitle('Scatter Plot', size=20, y=1.05)
plt.show()                                                  # Show the plot
```



Fitting a Linear Regression Model

```
In [16]: sns.regplot(x="YearsExperience", y="log_salary", data=df);
```



```
In [17]: import statsmodels.formula.api as smf

model = smf.ols("log_salary~YearsExperience", data = df).fit()
model.summary()
```

Out[17]:

OLS Regression Results

| | | | |
|--------------------------|------------------|----------------------------|----------|
| Dep. Variable: | log_salary | R-squared: | 0.932 |
| Model: | OLS | Adj. R-squared: | 0.930 |
| Method: | Least Squares | F-statistic: | 383.6 |
| Date: | Fri, 21 Oct 2022 | Prob (F-statistic): | 7.03e-18 |
| Time: | 21:32:08 | Log-Likelihood: | 28.183 |
| No. Observations: | 30 | AIC: | -52.37 |
| Df Residuals: | 28 | BIC: | -49.56 |
| Df Model: | 1 | | |

Covariance Type: nonrobust

| | coef | std err | t | P> t | [0.025 | 0.975] |
|------------------------|---------|---------|---------|-------|--------|--------|
| Intercept | 10.5074 | 0.038 | 273.327 | 0.000 | 10.429 | 10.586 |
| YearsExperience | 0.1255 | 0.006 | 19.585 | 0.000 | 0.112 | 0.139 |

| | | | |
|-----------------------|-------|--------------------------|-------|
| Omnibus: | 0.826 | Durbin-Watson: | 1.438 |
| Prob(Omnibus): | 0.661 | Jarque-Bera (JB): | 0.812 |
| Skew: | 0.187 | Prob(JB): | 0.666 |
| Kurtosis: | 2.286 | Cond. No. | 13.2 |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Summary:

That's one of the strong points of statsmodels Summary shows that beta values required for straight line for each feature as 'Years Experience' and 'constant'

Which feature is important for the Target is determined by P(t) i.e. hypothesis p-value

The feature Years Experience = p value(beta2) = 0.000 The feature Const = p value(beta1) = 0.000

Hence, we Reject null hypothesis in both cases. So, Years Experience and Constant both are important feature for the Target prediction

Intercept(Cosntant) of best fitted line is 10.5074 Feature(YearsExperience) coefficient is 0.1255

R-sqaured error value is closer to 1 that means the regression model covers most part of the variance of the values of the response variable and can be termed as a good model.

In [18]: `%matplotliblib` notebook


In [19]:

```
# Created a scatter plot
plt.scatter(x,y)

# Defined the estimated regression line, so we can plot it later
yhat = 0.1255*x + 10.5074 #predicted y

# Plotting the regression line against the independent variable (Years Experience)
fig = plt.plot(x,yhat, lw=4, c='red', label='Regression line')

# Label the axes
plt.xlabel('Years Experience', fontsize = 10)
plt.ylabel('Salary', fontsize = 10)
plt.show()
```



```
In [20]: #log values - predcited target value
yhat.head()
```

```
Out[20]: 0    10.64545
1    10.67055
2    10.69565
3    10.75840
4    10.78350
Name: YearsExperience, dtype: float64
```

```
In [21]: #Anti-log - actual y values
y = np.exp(yhat)
y.head()
```

```
Out[21]: 0    42001.054328
1    43068.622727
2    44163.326214
3    47023.370416
4    48218.594324
Name: YearsExperience, dtype: float64
```

```
In [22]: #Co-efficients values
#beta1 and bet0 values
model.params
```

```
Out[22]: Intercept      10.507402
YearsExperience    0.125453
dtype: float64
```

```
In [23]: #t and p-Values for intercept and Years Experience.
print(model.tvalues, '\n', model.pvalues)
```


| | |
|-----------------|--------------|
| Intercept | 273.327166 |
| YearsExperience | 19.584833 |
| dtype: float64 | |
| Intercept | 1.604634e-49 |
| YearsExperience | 7.027439e-18 |
| dtype: float64 | |